

Chapter 6: Moving Beyond Linearity

Fabian Peri

October 20, 2018

```
# Libraries
library(ISLR)
attach(Wage)

# Polynomial Regression and Step Functions

fit=lm(wage~poly(age,4),data=Wage)
coef(summary(fit))

##              Estimate Std. Error   t value    Pr(>|t|)
## (Intercept)   111.70361   0.7287409 153.283015 0.000000e+00
## poly(age, 4)1   447.06785  39.9147851  11.200558 1.484604e-28
## poly(age, 4)2  -478.31581  39.9147851 -11.983424 2.355831e-32
## poly(age, 4)3   125.52169  39.9147851   3.144742 1.678622e-03
## poly(age, 4)4  -77.91118  39.9147851  -1.951938 5.103865e-02

fit2=lm(wage~poly(age,4,raw=T),data=Wage)
coef(summary(fit2))

##              Estimate Std. Error   t value    Pr(>|t|)
## (Intercept)    -1.841542e+02  6.004038e+01 -3.067172 0.0021802539
## poly(age, 4, raw = T)1  2.124552e+01  5.886748e+00  3.609042 0.0003123618
## poly(age, 4, raw = T)2 -5.638593e-01  2.061083e-01 -2.735743 0.0062606446
## poly(age, 4, raw = T)3  6.810688e-03  3.065931e-03  2.221409 0.0263977518
## poly(age, 4, raw = T)4 -3.203830e-05  1.641359e-05 -1.951938 0.0510386498

fit2a=lm(wage~age+I(age^2)+I(age^3)+I(age^4),data=Wage)
coef(fit2a)

##      (Intercept)          age      I(age^2)      I(age^3)      I(age^4)
## -1.841542e+02  2.124552e+01 -5.638593e-01  6.810688e-03 -3.203830e-05

fit2b=lm(wage~cbind(age,age^2,age^3,age^4),data=Wage)
agelims=range(age)
age.grid=seq(from=agelims[1],to=agelims[2])
preds=predict(fit,newdata=list(age=age.grid),se=TRUE)
se.bands=cbind(preds$fit+2*preds$se.fit,preds$fit-2*preds$se.fit)
par(mfrow=c(1,2),mar=c(4.5,4.5,1,1),oma=c(0,0,4,0))
plot(age,wage,xlim=agelims,cex=.5,col="darkgrey")
title("Degree-4 Polynomial",outer=T)
lines(age.grid,preds$fit,lwd=2,col="blue")
matlines(age.grid,se.bands,lwd=1,col="blue",lty=3)
preds2=predict(fit2,newdata=list(age=age.grid),se=TRUE)
max(abs(preds$fit-preds2$fit))

## [1] 7.81597e-11

fit.1=lm(wage~age,data=Wage)
fit.2=lm(wage~poly(age,2),data=Wage)
fit.3=lm(wage~poly(age,3),data=Wage)
```

```
fit.4=lm(wage~poly(age,4),data=Wage)
fit.5=lm(wage~poly(age,5),data=Wage)
anova(fit.1,fit.2,fit.3,fit.4,fit.5)
```

```
## Analysis of Variance Table
##
```

```
## Model 1: wage ~ age
## Model 2: wage ~ poly(age, 2)
## Model 3: wage ~ poly(age, 3)
## Model 4: wage ~ poly(age, 4)
## Model 5: wage ~ poly(age, 5)
##   Res.Df    RSS Df Sum of Sq      F    Pr(>F)
## 1    2998 5022216
## 2    2997 4793430  1    228786 143.5931 < 2.2e-16 ***
## 3    2996 4777674  1     15756  9.8888 0.001679 **
## 4    2995 4771604  1      6070  3.8098 0.051046 .
## 5    2994 4770322  1      1283  0.8050 0.369682
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coef(summary(fit.5))
```

```
##              Estimate Std. Error    t value    Pr(>|t|)
## (Intercept)   111.70361   0.7287647  153.2780243 0.000000e+00
## poly(age, 5)1   447.06785  39.9160847   11.2001930 1.491111e-28
## poly(age, 5)2 -478.31581  39.9160847  -11.9830341 2.367734e-32
## poly(age, 5)3  125.52169  39.9160847   3.1446392 1.679213e-03
## poly(age, 5)4  -77.91118  39.9160847  -1.9518743 5.104623e-02
## poly(age, 5)5  -35.81289  39.9160847  -0.8972045 3.696820e-01
```

```
(-11.983)^2
```

```
## [1] 143.5923
```

```
fit.1=lm(wage~education+age,data=Wage)
fit.2=lm(wage~education+poly(age,2),data=Wage)
fit.3=lm(wage~education+poly(age,3),data=Wage)
anova(fit.1,fit.2,fit.3)
```

```
## Analysis of Variance Table
##
```

```
## Model 1: wage ~ education + age
## Model 2: wage ~ education + poly(age, 2)
## Model 3: wage ~ education + poly(age, 3)
##   Res.Df    RSS Df Sum of Sq      F    Pr(>F)
## 1    2994 3867992
## 2    2993 3725395  1    142597 114.6969 <2e-16 ***
## 3    2992 3719809  1      5587  4.4936 0.0341 *
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
fit=glm(I(wage>250)~poly(age,4),data=Wage,family=binomial)
preds=predict(fit,newdata=list(age=age.grid),se=T)
pfit=exp(preds$fit)/(1+exp(preds$fit))
se.bands.logit = cbind(preds$fit+2*preds$se.fit, preds$fit-2*preds$se.fit)
se.bands = exp(se.bands.logit)/(1+exp(se.bands.logit))
```

```

preds=predict(fit,newdata=list(age=age.grid),type="response",se=T)
plot(age,I(wage>250),xlim=agelims,type="n",ylim=c(0,.2))
points(jitter(age), I((wage>250)/5),cex=.5,pch="|",col="darkgrey")
lines(age.grid,pfit,lwd=2, col="blue")
matlines(age.grid,se.bands,lwd=1,col="blue",lty=3)

```

Degree-4 Polynomial



```
table(cut(age,4))
```

```
##
## (17.9,33.5]  (33.5,49]  (49,64.5] (64.5,80.1]
##           750      1399      779      72

```

```
fit=lm(wage~cut(age,4),data=Wage)
coef(summary(fit))

```

```
##
##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept)    94.158392    1.476069  63.789970 0.000000e+00
## cut(age, 4)(33.5,49]    24.053491    1.829431  13.148074 1.982315e-38
## cut(age, 4)(49,64.5]    23.664559    2.067958  11.443444 1.040750e-29
## cut(age, 4)(64.5,80.1]    7.640592    4.987424   1.531972 1.256350e-01

```

Splines

```
library(splines)
fit=lm(wage~bs(age,knots=c(25,40,60)),data=Wage)
pred=predict(fit,newdata=list(age=age.grid),se=T)
plot(age,wage,col="gray")

```

```

lines(age.grid,pred$fit,lwd=2)
lines(age.grid,pred$fit+2*pred$se,lty="dashed")
lines(age.grid,pred$fit-2*pred$se,lty="dashed")
dim(bs(age,knots=c(25,40,60)))

## [1] 3000    6

dim(bs(age,df=6))


## [1] 3000    6

attr(bs(age,df=6),"knots")

##    25%    50%    75%
## 33.75 42.00 51.00

fit2=lm(wage~ns(age,df=4),data=Wage)
pred2=predict(fit2,newdata=list(age=age.grid),se=T)
lines(age.grid, pred2$fit,col="red",lwd=2)

```



```

plot(age,wage,xlim=agelims,cex=.5,col="darkgrey")
title("Smoothing Spline")
fit=smooth.spline(age,wage,df=16)
fit2=smooth.spline(age,wage,cv=TRUE)

```

```

## Warning in smooth.spline(age, wage, cv = TRUE): cross-validation with non-
## unique 'x' values seems doubtful

```


```
fit2$df
```

```
## [1] 6.794596
```

```
lines(fit,col="red",lwd=2)
```

```
lines(fit2,col="blue",lwd=2)
```

```
legend("topright",legend=c("16 DF","6.8 DF"),col=c("red","blue"),lty=1,lwd=2,cex=.8)
```



```
plot(age,wage,xlim=agelims,cex=.5,col="darkgrey")
```

```
title("Local Regression")
```

```
fit=loess(wage~age,span=.2,data=Wage)
```


```
fit2=loess(wage~age,span=.5,data=Wage)
```

```
lines(age.grid,predict(fit,data.frame(age=age.grid)),col="red",lwd=2)
```

```
lines(age.grid,predict(fit2,data.frame(age=age.grid)),col="blue",lwd=2)
```

```
legend("topright",legend=c("Span=0.2","Span=0.5"),col=c("red","blue"),lty=1,lwd=2,cex=.8)
```

Local Regression



GAMs

```
gam1=lm(wage~ns(year,4)+ns(age,5)+education,data=Wage)
library(gam)
```

```
## Loading required package: foreach
```

```
## Loaded gam 1.16
```

```
gam.m3=gam(wage~s(year,4)+s(age,5)+education,data=Wage)
par(mfrow=c(1,3))
plot(gam.m3, se=TRUE,col="blue")
#plot.gam(gam1, se=TRUE, col="red")
gam.m1=gam(wage~s(age,5)+education,data=Wage)
gam.m2=gam(wage~year+s(age,5)+education,data=Wage)
anova(gam.m1,gam.m2,gam.m3,test="F")
```

```
## Analysis of Deviance Table
```

```
##
```

```
## Model 1: wage ~ s(age, 5) + education
```

```
## Model 2: wage ~ year + s(age, 5) + education
```

```
## Model 3: wage ~ s(year, 4) + s(age, 5) + education
```

```
##   Resid. Df Resid. Dev Df Deviance      F    Pr(>F)
```

```
## 1      2990      3711731
```

```
## 2      2989      3693842   1  17889.2 14.4771 0.0001447 ***
```

```
## 3      2986      3689770   3   4071.1  1.0982 0.3485661
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(gam.m3)
```

```
##
## Call: gam(formula = wage ~ s(year, 4) + s(age, 5) + education, data = Wage)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -119.43  -19.70   -3.33   14.17  213.48
##
## (Dispersion Parameter for gaussian family taken to be 1235.69)
##
##      Null Deviance: 5222086 on 2999 degrees of freedom
## Residual Deviance: 3689770 on 2986 degrees of freedom
## AIC: 29887.75
##
## Number of Local Scoring Iterations: 2
##
## Anova for Parametric Effects
##              Df Sum Sq Mean Sq F value    Pr(>F)
## s(year, 4)     1  27162   27162  21.981 2.877e-06 ***
## s(age, 5)       1 195338  195338 158.081 < 2.2e-16 ***
## education      4 1069726  267432 216.423 < 2.2e-16 ***
## Residuals    2986 3689770    1236
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##              Npar Df Npar F  Pr(F)
## (Intercept)
## s(year, 4)         3  1.086 0.3537
## s(age, 5)          4 32.380 <2e-16 ***
## education
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

preds=predict(gam.m2,newdata=Wage)
gam.lo=gam(wage~s(year,df=4)+lo(age,span=0.7)+education,data=Wage)
#plot.gam(gam.lo, se=TRUE, col="green")
gam.lo.i=gam(wage~lo(year,age,span=0.5)+education,data=Wage)


## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame,
## bf.maxit, : liv too small. (Discovered by lowesd)

## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame,
## bf.maxit, : lv too small. (Discovered by lowesd)


## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame,
## bf.maxit, : liv too small. (Discovered by lowesd)

## Warning in lo.wam(x, z, wz, fit$smooth, which, fit$smooth.frame,
## bf.maxit, : lv too small. (Discovered by lowesd)
```


```
library(akima)
```



```
plot(gam.lo.i)
gam.lr=gam(I(wage>250)~year+s(age,df=5)+education,family=binomial,data=Wage)
par(mfrow=c(1,3))
```

```
plot(gam.lm, se=T, col="green")
```



```
table(education,I(wage>250))
```

```
##
## education      FALSE TRUE
## 1. < HS Grad    268    0
## 2. HS Grad     966    5
## 3. Some College 643    7
## 4. College Grad 663   22
## 5. Advanced Degree 381  45
```

```
gam.lr.s=gam(I(wage>250)~year+s(age,df=5)+education,family=binomial,data=Wage,subset=(education!="1. < HS Grad"))
plot(gam.lr.s,se=T,col="green")
```

